

# Aritra Sengupta

## Curriculum Vitae

108 Naramore Ln.  
Los Gatos, CA-94032  
☎ (+1) 614-477-9964  
✉ aritra1686vit@gmail.com  
[link-to-webpage](#)  
US permanent resident

---

### INTERESTS

Program analysis (static and dynamic), parallel and distributed systems software, compilers, Responsible AI using PL and GenAI.

I enjoy designing and implementing analyses and runtimes across the software stack considering tradeoffs between performance and correctness.

I have a consistent track record of autonomously as well as collaboratively, delivering value, by delivering projects such as industry-scale software, product features with impact, and research papers and artifacts end to end. Experience in leading projects as a science lead working closely with product and engineering.

---

### EDUCATION

9/2011–5/2017 **Doctor of Philosophy (Ph.D.)**, *Department of Computer Sc. and Engg.*, The Ohio State University, Columbus, Ohio.

**Major:** Computer Science and Engg.

**Advisor:** Prof. Michael D. Bond.

**GPA:** 3.95/4

6/2004–6/2008 **Bachelor of Technology (B.Tech)**, *School of Computer Science and Engineering*, Vellore Institute of Technology University, Vellore, India.

**Major:** Computer Science and Engg.

**GPA:** 9.29/10

---

### SELECTED PUBLICATIONS

Subarno Banerjee, Siwei Cui, Michael Emmi, Antonio Filieri, Liana Hadarean, Peixuan Li, Linghui Luo, Goran Piskachev, Nicolás Rosner, **Sengupta, Aritra**, Omer Tripp, and Jingbo Wang. Compositional Taint Analysis for Large-Scale Policy Enforcement. In *ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*, 2023.

Michael Emmi, Liana Hadarean, Ranjit Jhala, Lee Pike, Nicolás Rosner, Martin Schäff, **Sengupta, Aritra**, and Willem Visser. RAPID: Checking API Usage for the Cloud in the Cloud. In *ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*, 2021.

Tanakorn Leesatapornwongsa, **Aritra Sengupta**, Masoud Saeida Ardekani, Gustavo Petri, and Cesar A. Stuardo. Transactuations—Where Transactions Meet the Physical World (extended version of USENIX ATC 2019 paper with formalized semantics). *ACM Trans. Comput. Syst. (TOCS)*, 36(4), May 2020.

**Aritra Sengupta**, Tanakorn Leesatapornwongsa, Masoud Saeida Ardekani, and Cesar Stuardo. Transactuation—Where Transaction Meets the Physical World. *Appeared at Usenix Annual Technical Conference (USENIX ATC, Awarded Best Paper)*, July 2019.

**Aritra Sengupta**, Man Cao, Michael D. Bond, and Milind Kulkarni. Legato: End-to-End Bounded Region Serializability Using Commodity Hardware Transactional Memory. In *ACM SIGPLAN International Symposium on Code Generation and Optimization (CGO)*, Feb 2017.

**Aritra Sengupta**, Man Cao, Michael D. Bond and Milind Kulkarni. Toward Efficient Strong Memory Model Support for the Java Platform via Hybrid Synchronization. In *ACM International Conference on Principles and Practices of Programming on the Java Platform (PPPJ'15)*, September 2015.

**Aritra Sengupta**, Swarnendu Biswas, Minjia Zhang, Michael D. Bond and Milind Kulkarni. EnfoRSer: Hybrid Static-Dynamic Analysis for Region Serializability. In *ACM Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, March 2015.

---

#### TALKS AND PRESENTATIONS

- Presented paper at *FSE 2021* (virtual)
- Presented paper at *USENIX ATC 2019*, Renton
- Presented paper at *CGO 2017*, Austin
- Presented paper at *ASPLOS 2015*, Istanbul
- Presented paper at *PPPJ 2015*, Florida
- Talks on PhD thesis at *Intel*, *Samsung Research*, *Mathworks* in 2017.
- Presented work at *ACM Student Research Competition*, *OOPSLA 2015*.
- Presented work at *ACM Student Research Competition*, *PLDI 2013*.

---

#### EXPERIENCE

7/2023–present **Senior Applied Scientist**, AWS AI, Santa Clara, California.

- High quality and secure code generation with Amazon Q leveraging machine learning and program analysis
- Pre-training data cleaning using program analysis, preference-tuning LLMs for code quality and security.

4/2022–06/2023 **Senior Applied Scientist**, AWS AI, Santa Clara, California.

- Designed and implemented a modular and scalable version of a program analysis tool with strong abstractions for soundness.
- Implemented scalability techniques to reduce scope of program analysis while retaining soundness and precision. Leveraged the analysis to launch a high-fidelity analysis as part of an external facing AWS service.
- Led publication of the above work. The paper to appear at FSE 2023.
- Developed techniques to analyze complex program constructs such as higher-order functions, that are typically hard to analyze precisely.

8/2019–04/2022 **Applied Scientist**, AWS, *Automated Reasoning Group* then AWS AI, Cupertino/Santa Clara, California.

- Checking/Enforcing correctness properties at scale on AWS services.
- Designed, built, and productionalized a Java program analysis tool to check properties on code at scale. The tool is currently deployed on analysis that scans every code review at Amazon and as part of multiple external facing AWS services.
- Led publication of the above work end to end. The paper appeared at FSE 2021.
- Designed, built, and productionalized a capability to provide evidence of properties checked that didn't result in detected vulnerabilities. The tool collected meaningful sequence of transitions that led the detector instance to terminate in a "good" state.

7/2017–8/2019 **Senior Research Engineer**, *Samsung Research America*, Mountain View, California.

- Built a reliable IoT application execution system that provides strong guarantees, e.g. keeping soft states/app states consistent with hard states/device states. A simple programming abstraction allows the IoT developer to express consistency requirements which is respected by our runtime. The runtime is built atop Azure serverless functions executing over Cosmos DB.
- Led publication of the above work end to end. The paper appeared at USENIX ATC, 2019. Awarded Best Paper Award!
- Proposed a transactional framework to provide low-latency and serializable execution of smart-home applications. Development of a static analysis framework to understand the interaction of smart-home applications operating on shared software states and physical actuators. Transactions in the proposed framework read speculative soft states from the hub instead of physical sensors and actuators to derive a consistent order before issuing actual physical actuations on the devices.

9/2011–5/2017 **Graduate Research Associate**, *Computer Science and Engineering Department*, *Ohio State University*, Columbus, Ohio.

**Mentor:** Michael D. Bond

- Designed and implemented large concurrent systems considering tradeoffs between performance and correctness across the software stack.
- Built analysis on top of a default runtime and extracted performance considering several avenues of optimization in design and implementation [ASPLOS 2015, PPPJ 2015, CGO 2017, PPOPP 2016, OOPSLA 2013].
- Designed and implemented low-overhead techniques using hybrid static—dynamic analysis to enforce bounded region serializability. The techniques require compiler transformations, code generation, implementation of log-based and register-based transactional memory, efficient compiler analyses, etc.[ASPLOS 2015, PPPJ 2015, CGO 2017].
- Developed static analysis to eliminate redundant instrumentation, demarcate regions of code, generate dependence graphs for efficient code generation during JIT compilation [ASPLOS 2015, OOPSLA 2013].
- Designed and implemented dynamic analysis to detect shared memory bugs including sequential consistency violations and atomicity violations [ISMM 2016, ASPLOS 2015].
- Designed and Implemented hybrid synchronization techniques (per-access lock, per-region lock) to enable low-overhead enforcement of a strong memory model. Implementation of pessimistic static locks and optimistic dynamic locks [PPPJ 2015].
- Leveraged hardware transactional memory (HTM), conditional code generation, and control-theory algorithms to enforce a strong, always-on and practical memory consistency model [CGO 2017].
- Developed efficient and practical instrumentation schemes, implementing and modifying several parts of a JVM including baseline compiler, optimizing compiler, Java locking protocols, garbage collection features etc. [ASPLOS 2015, PPPJ 2015, ISMM 2016, PPOPP 2016, OOPSLA 2013, PLDI 2014, CGO 2017].

5/2016–8/2016

**Summer Research Intern**, *Microsoft Research*, Redmond, Seattle, USA.

**Mentors:** Kathryn S. McKinley, Ricardo Bianchini, Sameh Elnikety, Yuxiong He.

- Research on designing dynamic and adaptive scheduling techniques for latency sensitive server systems. Design and implementation of an online technique to detect memory interference amongst co-scheduled requests, using hardware performance counters, to automatically manage the server's scheduling, concurrency and intra-request parallelism in order to increase server utilization and reduce tail-latency.

---

#### SERVICE

- Program committee member for HiPC 2024
- Program committee member for ASPLOS 2024
- Program committee member for ASPLOS 2023
- Program committee member for OOPSLA 2021.
- Session chair, OOPSLA 2020.
- Program committee member for OOPSLA 2020.
- Reviewer CONCUR 2020.
- Reviewer for journal: IEEE Transactions on Software Engineering(TSE).
- Program committee member for IPDPS 2019.
- External review committee (ERC) member for ASPLOS 2019.
- Reviewer for PACT 2018.
- Reviewer for journal: ACM Transactions on Architecture and Code Optimization(TACO).
- Reviewer for journal: IEEE Transactions on Parallel and Distributed Systems (TPDS).
- Reviewer for journal: Concurrency and Computation: Practice and Experience.
- Member of EuroSys 2018 shadow program committee.
- Member of ASPLOS 2018 shadow program committee.
- Artifact evaluation committee member for PPOPP 2017/CGO 2017.
- Artifact evaluation committee member for PPOPP 2016/CGO 2016.
- Subreviewer for PLDI 2017.
- Subreviewer for PACT 2017.

---

#### MENTORING EXPERIENCE

- Mentored summer 2022 research intern at AWS CodeGuru—Siwei Cui, Ph.D. student at Texas A&M university.
- Mentored summer 2021 research intern at AWS CodeGuru—Ali Nowraiz Khan, Ph.D. student at U.C. Riverside.
- Mentored summer 2020 research intern at AWS CodeGuru—Ali Nowraiz Khan, Ph.D. student at U.C. Riverside.

- Mentored summer 2018 research intern at Samsung Research—Cesar S. Stuardo, Ph.D. student at U.Chicago.

---

#### HONORS AND AWARDS

- 7/2019 Awarded the Best Paper Award for Transactuation at USENIX ATC 2019.
- 2/2017 Awarded NSF travel grant for presenting at CGO 2017 in Austin, USA
- 3/2015 Awarded NSF travel grant for presenting at ASPLOS 2015 in Istanbul, Turkey.
- 2011-2012 Awarded the prestigious “University Fellowship”, graduate student fellowsip, Ohio State University.
- 11/2008 Secured “Initial Learning Program Top Performer” award in the first phase of training at Tata Consultancy Services, India.
- 2006-2008 Awarded merit certificate and scholarship in three consecutive years 2006, 2007, 2008 for academic performance at Vellore Institute of Technology University, Vellore, India.
- 2004 Awarded merit certificate for academic performance in Physics, Chemistry and Mathematics in Indian School Certificate Examination.

---

#### TECHNICAL SKILLS

- Java, Python, Bash
- Git
- Asynchronous programming using Node.js—Javascript.
- Tree-sitter
- Program analysis, compilers, concurrency, and distributed systems
- Parallel programming libraries, Performance engineering
- Responsible AI using PL and GenAI

---

#### PATENTS

Please refer webpage for granted and filed patents.